

Linux, Docker & AWS Deployment Guide

Table of Contents

1. [Linux Fundamentals](#)
 2. [Docker Essentials](#)
 3. [AWS Basics](#)
 4. [Deploying Flask App to AWS](#)
 5. [Dockerizing Flask Applications](#)
 6. [Deploying Docker Containers on AWS EC2](#)
 7. [Advanced Deployment Strategies](#)
-

Linux Fundamentals

Essential Commands

File and Directory Operations

bash

Navigation

```
pwd                # Print working directory
ls -la             # List files with details
cd /path/to/directory # Change directory
cd ~               # Go to home directory
cd ..              # Go up one directory
```

File operations

```
touch filename.txt # Create empty file
mkdir dirname      # Create directory
mkdir -p path/to/dir # Create nested directories
cp source dest     # Copy files
mv source dest     # Move/rename files
rm filename        # Remove file
rm -rf dirname     # Remove directory recursively
ln -s target link  # Create symbolic link
```

File viewing and editing

```
cat filename      # Display file content
less filename     # View file with pagination
head -10 filename # Show first 10 lines
tail -10 filename # Show last 10 lines
tail -f filename  # Follow file changes (logs)
nano filename     # Simple text editor
vim filename      # Advanced text editor
```

File Permissions

bash

Understanding permissions (rwx for user, group, others)

`ls -l filename`

-rw-r--r-- 1 user group 1024 Jan 1 12:00 filename

||| ||| |||

||| ||| ||+-- Others permissions

||| |||+---- Group permissions

|||+----- User permissions

+------ File type

Changing permissions

`chmod 755 filename` *# rwxr-xr-x*

`chmod +x script.sh` *# Add execute permission*

`chmod u+w filename` *# Add write permission for user*

`chown user:group file` *# Change ownership*

Common permission patterns

`chmod 644 file.txt` *# Read/write for owner, read for others*

`chmod 755 script.sh` *# Execute permissions for scripts*

`chmod 600 private.key` *# Read/write for owner only*

Process Management

bash

Process information

`ps aux` *# Show all running processes*
`ps aux | grep python` *# Find Python processes*
`top` *# Real-time process monitor*
`htop` *# Enhanced process monitor*
`pgrep -f flask` *# Find processes by name*

Process control

`kill PID` *# Terminate process by ID*
`kill -9 PID` *# Force kill process*
`killall python` *# Kill all Python processes*
`jobs` *# Show background jobs*
`bg` *# Put job in background*
`fg` *# Bring job to foreground*
`nohup command &` *# Run command that survives Logout*

Service management (systemd)

`sudo systemctl start service-name`
`sudo systemctl stop service-name`
`sudo systemctl restart service-name`
`sudo systemctl enable service-name` *# Auto-start on boot*
`sudo systemctl status service-name`

Network and System Information

bash

Network

`wget https://example.com/file.zip` *# Download files*
`curl -X GET https://api.example.com` *# Make HTTP requests*
`netstat -tulpn` *# Show network connections*
`ss -tulpn` *# Modern alternative to netstat*
`ping google.com` *# Test connectivity*
`traceroute google.com` *# Trace network path*

System information

`uname -a` *# System information*
`df -h` *# Disk usage*
`du -sh *` *# Directory sizes*
`free -h` *# Memory usage*
`uptime` *# System uptime and Load*
`whoami` *# Current user*
`id` *# User and group IDs*

Package Management (Ubuntu/Debian)

```
bash

# APT package manager
sudo apt update           # Update package list
sudo apt upgrade          # Upgrade installed packages
sudo apt install package-name # Install package
sudo apt remove package-name # Remove package
sudo apt autoremove       # Remove unused packages
apt search keyword        # Search for packages
apt show package-name     # Show package information

# Adding repositories
sudo add-apt-repository ppa:repo-name
sudo apt-key add key-file

# Snap packages
sudo snap install package-name
sudo snap list
sudo snap remove package-name
```

Environment Variables and Shell

```
bash

# Environment variables
echo $PATH           # Show PATH variable
export VAR_NAME="value" # Set environment variable
export PATH=$PATH:/new/path # Add to PATH
env                  # Show all environment variables
unset VAR_NAME       # Remove environment variable

# Shell configuration
nano ~/.bashrc       # Edit bash configuration
source ~/.bashrc      # Reload configuration
alias ll='ls -la'     # Create command alias
history               # Show command history
which python3         # Find command location
type python3          # Show command type and location
```

Docker Essentials

Docker Concepts

Core Components

- **Image:** Read-only template for creating containers
- **Container:** Running instance of an image
- **Dockerfile:** Text file with instructions to build an image
- **Registry:** Storage for Docker images (Docker Hub, AWS ECR)
- **Volume:** Persistent data storage for containers

Basic Docker Commands

Image Management

```
bash
```

```
# Pull images from registry
```

```
docker pull ubuntu:20.04
```

```
docker pull python:3.9-slim
```

```
docker pull nginx:alpine
```

```
# List images
```

```
docker images
```

```
docker image ls
```

```
# Remove images
```

```
docker rmi image-name:tag
```

```
docker rmi image-id
```

```
docker image prune # Remove unused images
```

```
# Build image from Dockerfile
```

```
docker build -t my-app:latest .
```

```
docker build -t my-app:v1.0 -f Dockerfile.prod .
```

```
# Tag images
```

```
docker tag my-app:latest my-registry.com/my-app:latest
```

```
# Push to registry
```

```
docker push my-registry.com/my-app:latest
```

Container Management

bash

Run containers

```
docker run hello-world           # Simple run
docker run -d nginx              # Run in background (detached)
docker run -p 8080:80 nginx      # Port mapping
docker run -v /host/path:/container/path ubuntu # Volume mounting
docker run -e ENV_VAR=value ubuntu # Environment variables
docker run --name my-container nginx # Named container
docker run -it ubuntu bash       # Interactive terminal
```

List containers

```
docker ps           # Running containers
docker ps -a        # ALL containers (including stopped)
```

Container operations

```
docker start container-name # Start stopped container
docker stop container-name  # Stop running container
docker restart container-name # Restart container
docker pause container-name # Pause container
docker unpause container-name # Unpause container
```

Execute commands in running container

```
docker exec -it container-name bash
docker exec container-name ls /app
```

View Logs

```
docker logs container-name
docker logs -f container-name # Follow Logs
```

Remove containers

```
docker rm container-name
docker rm -f container-name # Force remove running container
docker container prune      # Remove all stopped containers
```

Docker Compose

yaml

```
# docker-compose.yml
```

```
version: '3.8'
```

```
services:
```

```
  web:
```

```
    build: .
```

```
    ports:
```

```
      - "5000:5000"
```

```
    environment:
```

```
      - FLASK_ENV=development
```

```
    volumes:
```

```
      - ./app
```

```
    depends_on:
```

```
      - db
```

```
  db:
```

```
    image: postgres:13
```

```
    environment:
```

```
      POSTGRES_DB: myapp
```

```
      POSTGRES_USER: user
```

```
      POSTGRES_PASSWORD: password
```

```
    volumes:
```

```
      - postgres_data:/var/lib/postgresql/data
```

```
    ports:
```

```
      - "5432:5432"
```

```
volumes:
```

```
  postgres_data:
```

bash

```
# Docker Compose commands
```

```
docker-compose up
```

```
docker-compose up -d
```

```
docker-compose down
```

```
docker-compose build
```

```
docker-compose logs web
```

```
docker-compose exec web bash
```

```
# Start services
```

```
# Start in background
```

```
# Stop and remove services
```

```
# Build services
```

```
# View service logs
```

```
# Execute command in service
```

Creating Dockerfiles

Basic Dockerfile Structure

dockerfile

Use official Python runtime as base image

FROM python:3.9-slim

Set working directory in container

WORKDIR /app

Copy requirements first (for better caching)

COPY requirements.txt .

Install dependencies

RUN pip install --no-cache-dir -r requirements.txt

Copy application code

COPY . .

Expose port

EXPOSE 5000

Set environment variables

ENV FLASK_APP=app.py

ENV FLASK_ENV=production

Create non-root user for security

RUN adduser --disabled-password --gecos '' appuser

USER appuser

Command to run application

CMD ["python", "app.py"]

Multi-stage Dockerfile

dockerfile

Build stage

FROM python:3.9 as builder

WORKDIR /app

COPY requirements.txt .

RUN pip install --user -r requirements.txt

Production stage

FROM python:3.9-slim

WORKDIR /app

Copy installed packages from builder stage

COPY --from=builder /root/.local /root/.local

Copy application

COPY . .

Make sure scripts in .local are usable

ENV PATH=/root/.local/bin:\$PATH

EXPOSE 5000

CMD ["python", "app.py"]

AWS Basics

Core AWS Services

Compute Services

- **EC2**: Virtual servers in the cloud
- **Lambda**: Serverless compute functions
- **ECS**: Container orchestration service
- **Fargate**: Serverless containers

Storage Services

- **S3**: Object storage service
- **EBS**: Block storage for EC2
- **EFS**: Network file system

Database Services

- **RDS**: Managed relational databases
- **DynamoDB**: NoSQL database
- **ElastiCache**: In-memory caching

Networking

- **VPC**: Virtual Private Cloud
- **Route 53**: DNS service
- **CloudFront**: Content Delivery Network
- **Load Balancer**: Distribute traffic

AWS CLI Setup

Installation and Configuration

```
bash
```

```
# Install AWS CLI
```

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
```

```
unzip awscliv2.zip
```

```
sudo ./aws/install
```

```
# Configure AWS CLI
```

```
aws configure
```

```
# AWS Access Key ID: YOUR_ACCESS_KEY
```

```
# AWS Secret Access Key: YOUR_SECRET_KEY
```

```
# Default region name: us-east-1
```

```
# Default output format: json
```

```
# Verify configuration
```

```
aws sts get-caller-identity
```

```
aws s3 ls
```

Common AWS CLI Commands

```
bash
```

EC2 operations

```
aws ec2 describe-instances
```

```
aws ec2 start-instances --instance-ids i-1234567890abcdef0
```

```
aws ec2 stop-instances --instance-ids i-1234567890abcdef0
```

```
aws ec2 create-security-group --group-name my-sg --description "My security group"
```

S3 operations

```
aws s3 ls
```

```
aws s3 cp file.txt s3://my-bucket/
```

```
aws s3 sync ./local-folder s3://my-bucket/remote-folder/
```

```
aws s3 mb s3://my-new-bucket
```

IAM operations

```
aws iam list-users
```

```
aws iam create-user --user-name myuser
```

```
aws iam attach-user-policy --user-name myuser --policy-arn arn:aws:iam::aws:policy/ReadOnlyAccess
```

Deploying Flask App to AWS

Method 1: Direct Deployment on EC2

Step 1: Launch EC2 Instance

```
bash
```

Create key pair

```
aws ec2 create-key-pair --key-name my-flask-key --query 'KeyMaterial' --output text > my-flask-key.pem  
chmod 400 my-flask-key.pem
```

Launch instance

```
aws ec2 run-instances \
  --image-id ami-0c02fb55956c7d316 \
  --count 1 \
  --instance-type t2.micro \
  --key-name my-flask-key \
  --security-groups flask-app-sg
```

Step 2: Connect to Instance

```
bash

# Get instance public IP
aws ec2 describe-instances --query 'Reservations[*].Instances[*].PublicIpAddress' --output text

# Connect via SSH
ssh -i my-flask-key.pem ubuntu@YOUR_INSTANCE_IP
```

Step 3: Setup Environment on EC2

```
bash

# Update system
sudo apt update && sudo apt upgrade -y

# Install Python and dependencies
sudo apt install -y python3 python3-pip python3-venv nginx

# Create application directory
sudo mkdir -p /var/www/flask-app
sudo chown ubuntu:ubuntu /var/www/flask-app
cd /var/www/flask-app

# Create virtual environment
python3 -m venv venv
source venv/bin/activate

# Clone your application (or upload files)
git clone https://github.com/yourusername/your-flask-app.git .
# OR upload files using scp:
# scp -i my-flask-key.pem -r ./my-flask-app ubuntu@YOUR_INSTANCE_IP:/var/www/flask-app/

# Install dependencies
pip install -r requirements.txt
pip install gunicorn
```

Step 4: Configure Gunicorn

bash

Create Gunicorn configuration

```
cat > gunicorn.conf.py << EOF
```

```
bind = "127.0.0.1:5000"
```

```
workers = 2
```

```
worker_class = "sync"
```

```
worker_connections = 1000
```

```
max_requests = 1000
```

```
max_requests_jitter = 100
```

```
timeout = 30
```

```
keepalive = 5
```

```
preload_app = True
```

```
EOF
```

Create systemd service file

```
sudo tee /etc/systemd/system/flask-app.service << EOF
```

```
[Unit]
```

```
Description=Gunicorn instance to serve Flask App
```

```
After=network.target
```

```
[Service]
```

```
User=ubuntu
```

```
Group=www-data
```

```
WorkingDirectory=/var/www/flask-app
```

```
Environment="PATH=/var/www/flask-app/venv/bin"
```

```
ExecStart=/var/www/flask-app/venv/bin/gunicorn --config gunicorn.conf.py app:app
```

```
ExecReload=/bin/kill -s HUP \${MAINPID}
```

```
Restart=always
```

```
[Install]
```

```
WantedBy=multi-user.target
```

```
EOF
```

Start and enable service

```
sudo systemctl daemon-reload
```

```
sudo systemctl start flask-app
```

```
sudo systemctl enable flask-app
```

```
sudo systemctl status flask-app
```

Step 5: Configure Nginx

bash

Create Nginx configuration

```
sudo tee /etc/nginx/sites-available/flask-app << EOF
```

```
server {  
    listen 80;  
    server_name YOUR_DOMAIN_OR_IP;  
  
    location / {  
        proxy_pass http://127.0.0.1:5000;  
        proxy_set_header Host $host;  
        proxy_set_header X-Real-IP $remote_addr;  
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;  
        proxy_set_header X-Forwarded-Proto $scheme;  
    }  
  
    # Optional: Serve static files directly  
    location /static {  
        alias /var/www/flask-app/static;  
        expires 1y;  
        add_header Cache-Control "public, immutable";  
    }  
}  
EOF
```

Enable site and restart Nginx

```
sudo ln -s /etc/nginx/sites-available/flask-app /etc/nginx/sites-enabled/
```

```
sudo nginx -t
```

```
sudo systemctl restart nginx
```

Step 6: Configure Security Group

bash

Create security group

```
aws ec2 create-security-group \  
  --group-name flask-app-sg \  
  --description "Security group for Flask application"
```

Allow HTTP traffic

```
aws ec2 authorize-security-group-ingress \  
  --group-name flask-app-sg \  
  --protocol tcp \  
  --port 80 \  
  --cidr 0.0.0.0/0
```

Allow HTTPS traffic

```
aws ec2 authorize-security-group-ingress \  
  --group-name flask-app-sg \  
  --protocol tcp \  
  --port 443 \  
  --cidr 0.0.0.0/0
```

Allow SSH access

```
aws ec2 authorize-security-group-ingress \  
  --group-name flask-app-sg \  
  --protocol tcp \  
  --port 22 \  
  --cidr 0.0.0.0/0
```

Method 2: Using AWS Elastic Beanstalk

Step 1: Prepare Application


```
bash
```

```
# Create application.py (Elastic Beanstalk expects this name)
```

```
cp app.py application.py
```

```
# Create requirements.txt
```

```
pip freeze > requirements.txt
```

```
# Create .ebextensions/python.config
```

```
mkdir .ebextensions
```

```
cat > .ebextensions/python.config << EOF
```

```
option_settings:
```

```
  aws:elasticbeanstalk:container:python:
```

```
    WSGIPath: application:app
```

```
  aws:elasticbeanstalk:environment:proxy:staticfiles:
```

```
    /static: static
```

```
EOF
```

Step 2: Deploy with EB CLI

```
bash
```

```
# Install EB CLI
```

```
pip install awsebcli
```

```
# Initialize Elastic Beanstalk application
```

```
eb init flask-app --region us-east-1 --platform python-3.9
```

```
# Create environment and deploy
```

```
eb create flask-app-env
```

```
# Deploy updates
```

```
eb deploy
```

```
# Open application in browser
```

```
eb open
```

```
# View logs
```

```
eb logs
```

```
# Terminate environment (cleanup)
```

```
eb terminate flask-app-env
```

Dockerizing Flask Applications

Complete Flask Application Structure

```
flask-docker-app/  
├─ app.py  
├─ requirements.txt  
├─ Dockerfile  
├─ docker-compose.yml  
├─ .dockerignore  
├─ nginx/  
│   └─ nginx.conf  
└─ static/  
    └─ styles.css
```

Sample Flask Application


```

# app.py
from flask import Flask, jsonify, request
import os
import redis
from datetime import datetime

app = Flask(__name__)

# Redis connection (optional)
try:
    redis_client = redis.Redis(
        host=os.environ.get('REDIS_HOST', 'localhost'),
        port=int(os.environ.get('REDIS_PORT', 6379)),
        decode_responses=True
    )
except:
    redis_client = None

@app.route('/')
def home():
    return jsonify({
        "message": "Flask Docker App",
        "timestamp": datetime.now().isoformat(),
        "environment": os.environ.get('FLASK_ENV', 'production')
    })

@app.route('/health')
def health():
    return jsonify({"status": "healthy"}), 200

@app.route('/api/counter', methods=['GET', 'POST'])
def counter():
    if not redis_client:
        return jsonify({"error": "Redis not available"}), 503

    if request.method == 'POST':
        count = redis_client.incr('counter')
        return jsonify({"count": count}), 201
    else:
        count = redis_client.get('counter') or 0
        return jsonify({"count": int(count)})

if __name__ == '__main__':
    app.run(
        host='0.0.0.0',
        port=int(os.environ.get('PORT', 5000)),

```

```
        debug=os.environ.get('FLASK_ENV') == 'development'  
    )
```

Production Dockerfile

dockerfile

Dockerfile

FROM python:3.9-slim

Set environment variables

ENV PYTHONDONTWRITEBYTECODE=1

ENV PYTHONUNBUFFERED=1

ENV FLASK_APP=app.py

ENV FLASK_ENV=production

Set work directory

WORKDIR /app

Install system dependencies

```
RUN apt-get update \  
    && apt-get install -y --no-install-recommends \  
        gcc \  
    && rm -rf /var/lib/apt/lists/*
```

Install Python dependencies

COPY requirements.txt .

RUN pip install --no-cache-dir -r requirements.txt

Copy project

COPY . .

Create non-root user

```
RUN adduser --disabled-password --gecos '' appuser \  
    && chown -R appuser:appuser /app
```

USER appuser

Expose port

EXPOSE 5000

Health check

```
HEALTHCHECK --interval=30s --timeout=10s --start-period=5s --retries=3 \  
    CMD curl -f http://localhost:5000/health || exit 1
```

Run application

CMD ["gunicorn", "--bind", "0.0.0.0:5000", "--workers", "2", "app:app"]

Docker Compose for Development

yaml

```
# docker-compose.yml
version: '3.8'

services:
  web:
    build: .
    ports:
      - "5000:5000"
    environment:
      - FLASK_ENV=development
      - REDIS_HOST=redis
      - REDIS_PORT=6379
    volumes:
      - ./app
    depends_on:
      - redis
    restart: unless-stopped

  redis:
    image: redis:7-alpine
    ports:
      - "6379:6379"
    volumes:
      - redis_data:/data
    restart: unless-stopped

  nginx:
    image: nginx:alpine
    ports:
      - "80:80"
    volumes:
      - ./nginx/nginx.conf:/etc/nginx/nginx.conf:ro
    depends_on:
      - web
    restart: unless-stopped

volumes:
  redis_data:
```

Production Docker Compose


```
# docker-compose.prod.yml
version: '3.8'

services:
  web:
    build:
      context: .
      dockerfile: Dockerfile
    environment:
      - FLASK_ENV=production
      - REDIS_HOST=redis
      - DATABASE_URL=postgresql://user:pass@db:5432/myapp
    depends_on:
      - redis
      - db
    restart: unless-stopped
    networks:
      - app-network

  nginx:
    image: nginx:alpine
    ports:
      - "80:80"
      - "443:443"
    volumes:
      - ./nginx/nginx.conf:/etc/nginx/nginx.conf:ro
      - ./nginx/ssl:/etc/nginx/ssl:ro
    depends_on:
      - web
    restart: unless-stopped
    networks:
      - app-network

  redis:
    image: redis:7-alpine
    volumes:
      - redis_data:/data
    restart: unless-stopped
    networks:
      - app-network

  db:
    image: postgres:13
    environment:
      POSTGRES_DB: myapp
      POSTGRES_USER: user
```



```
    POSTGRES_PASSWORD: password
volumes:
  - postgres_data:/var/lib/postgresql/data
restart: unless-stopped
networks:
  - app-network
```

```
volumes:
  redis_data:
  postgres_data:
```

```
networks:
  app-network:
    driver: bridge
```

Nginx Configuration

nginx

```
# nginx/nginx.conf
events {
    worker_connections 1024;
}

http {
    upstream app {
        server web:5000;
    }

    server {
        listen 80;
        server_name localhost;

        client_max_body_size 10M;

        location / {
            proxy_pass http://app;
            proxy_set_header Host $host;
            proxy_set_header X-Real-IP $remote_addr;
            proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
            proxy_set_header X-Forwarded-Proto $scheme;
            proxy_redirect off;
        }

        location /static/ {
            alias /app/static/;
            expires 1y;
            add_header Cache-Control "public, immutable";
        }
    }
}
```

.dockerignore

```
# .dockerignore
.git
.gitignore
README.md
Dockerfile
docker-compose*.yaml
.env
.venv
venv/
__pycache__/
*.pyc
*.pyo
*.pyd
.Python
.pytest_cache
.coverage
htmlcov/
.tox/
.cache
nosetests.xml
coverage.xml
```

Deploying Docker Containers on AWS EC2

Method 1: Docker on EC2

Step 1: Launch EC2 Instance with Docker

```
bash
```

```
# Create user data script for Docker installation
```

```
cat > user-data.sh << 'EOF'
```

```
#!/bin/bash
```

```
yum update -y
```

```
amazon-linux-extras install docker
```

```
service docker start
```

```
usermod -a -G docker ec2-user
```

```
systemctl enable docker
```

```
# Install Docker Compose
```

```
curl -L "https://github.com/docker/compose/releases/latest/download/docker-compose-$(uname -s)-
```

```
chmod +x /usr/local/bin/docker-compose
```

```
ln -s /usr/local/bin/docker-compose /usr/bin/docker-compose
```

```
EOF
```

```
# Launch instance with user data
```

```
aws ec2 run-instances \
```

```
--image-id ami-0c02fb55956c7d316 \
```

```
--count 1 \
```

```
--instance-type t3.small \
```

```
--key-name my-flask-key \
```

```
--security-group-ids sg-xxxxxxx \
```

```
--user-data file://user-data.sh \
```

```
--tag-specifications 'ResourceType=instance,Tags=[{Key=Name,Value=Flask-Docker-App}]'
```

Step 2: Deploy Application

```
bash
```

```
# Connect to instance
```

```
ssh -i my-flask-key.pem ec2-user@YOUR_INSTANCE_IP
```

```
# Clone repository
```

```
git clone https://github.com/yourusername/flask-docker-app.git
```

```
cd flask-docker-app
```

```
# Build and run with Docker Compose
```

```
docker-compose -f docker-compose.prod.yml up -d
```

```
# Check status
```

```
docker-compose -f docker-compose.prod.yml ps
```

```
docker-compose -f docker-compose.prod.yml logs
```

Method 2: Using Amazon ECR (Elastic Container Registry)

Step 1: Create ECR Repository


```
bash

# Create repository
aws ecr create-repository --repository-name flask-app

# Get Login token
aws ecr get-login-password --region us-east-1 | docker login --username AWS --password-stdin 123456789012.dkr.ecr.us-east-1.amazonaws.com

# Build and tag image
docker build -t flask-app .
docker tag flask-app:latest 123456789012.dkr.ecr.us-east-1.amazonaws.com/flask-app:latest

# Push image
docker push 123456789012.dkr.ecr.us-east-1.amazonaws.com/flask-app:latest
```




Step 2: Deploy from ECR

```
bash

# On EC2 instance, pull and run image
aws ecr get-login-password --region us-east-1 | docker login --username AWS --password-stdin 123456789012.dkr.ecr.us-east-1.amazonaws.com

docker pull 123456789012.dkr.ecr.us-east-1.amazonaws.com/flask-app:latest
docker run -d -p 80:5000 --name flask-app 123456789012.dkr.ecr.us-east-1.amazonaws.com/flask-app:latest
```



Method 3: Using Amazon ECS (Elastic Container Service)

Step 1: Create Task Definition

json

```
{
  "family": "flask-app-task",
  "networkMode": "awsvpc",
  "requiresCompatibilities": ["FARGATE"],
  "cpu": "256",
  "memory": "512",
  "executionRoleArn": "arn:aws:iam::123456789012:role/ecsTaskExecutionRole",
  "containerDefinitions": [
    {
      "name": "flask-app",
      "image": "123456789012.dkr.ecr.us-east-1.amazonaws.com/flask-app:latest",
      "portMappings": [
        {
          "containerPort": 5000,
          "protocol": "tcp"
        }
      ],
      "essential": true,
      "logConfiguration": {
        "logDriver": "awslogs",
        "options": {
          "awslogs-group": "/ecs/flask-app",
          "awslogs-region": "us-east-1",
          "awslogs-stream-prefix": "ecs"
        }
      },
      "environment": [
        {
          "name": "FLASK_ENV",
          "value": "production"
        }
      ]
    }
  ]
}
```

Step 2: Create ECS Cluster and Service

bash

Create cluster

```
aws ecs create-cluster --cluster-name flask-app-cluster
```

Register task definition

```
aws ecs register-task-definition --cli-input-json file://task-definition.json
```

Create service

```
aws ecs create-service \  
  --cluster flask-app-cluster \  
  --service-name flask-app-service \  
  --task-definition flask-app-task:1 \  
  --desired-count 2 \  
  --launch-type FARGATE \  
  --network-configuration "awsVpcConfiguration={subnets=[subnet-12345678],securityGroups=[sg-
```



Advanced Deployment Strategies

Auto Scaling with Application Load Balancer

Step 1: Create Load Balancer

bash

Create Application Load Balancer

```
aws elbv2 create-load-balancer \  
  --name flask-app-alb \  
  --subnets subnet-12345678 subnet-87654321 \  
  --security-groups sg-12345678
```

Create target group

```
aws elbv2 create-target-group \  
  --name flask-app-targets \  
  --protocol HTTP \  
  --port 80 \  
  --vpc-id vpc-12345678 \  
  --health-check-path /health
```

Create Listener

```
aws elbv2 create-listener \  
  --load-balancer-arn arn:aws:elasticloadbalancing:us-east-1:123456789012:loadbalancer/app/fl  
  --protocol HTTP \  
  --port 80 \  
  --default-actions Type=forward,TargetGroupArn=arn:aws:elasticloadbalancing:us-east-1:123456
```



Step 2: Auto Scaling Group

bash

Create Launch template

```
aws ec2 create-launch-template \  
  --launch-template-name flask-app-template \  
  --launch-template-data '{  
    "ImageId": "ami-0c02fb55956c7d316",  
    "InstanceType": "t3.micro",  
    "KeyName": "my-flask-key",  
    "SecurityGroupIds": ["sg-12345678"],  
    "UserData": "'$(base64 -w 0 user-data.sh)'"  
  }'
```

Create Auto Scaling Group

```
aws autoscaling create-auto-scaling-group \  
  --auto-scaling-group-name flask-app-asg \  
  --launch-template LaunchTemplateName=flask-app-template,Version=1 \  
  --min-size 1 \  
  --max-size 5 \  
  --desired-capacity 2 \  
  --target-group-arns arn:aws:elasticloadbalancing:us-east-1:123456789012:targetgroup/flask-a  
  --vpc-zone-identifier "subnet-12345678,subnet-87654321"
```

CI/CD Pipeline with GitHub Actions


```
# .github/workflows/deploy.yml
name: Deploy Flask App

on:
  push:
    branches: [main]
  pull_request:
    branches: [main]

env:
  AWS_REGION: us-east-1
  ECR_REPOSITORY: flask-app
  ECS_SERVICE: flask-app-service
  ECS_CLUSTER: flask-app-cluster
  ECS_TASK_DEFINITION: task-definition.json

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3

      - name: Set up Python
        uses: actions/setup-python@v3
        with:
          python-version: '3.9'

      - name: Install dependencies
        run: |
          python -m pip install --upgrade pip
          pip install -r requirements.txt
          pip install pytest pytest-cov

      - name: Run tests
        run: |
          pytest tests/ --cov=app --cov-report=xml

      - name: Upload coverage to Codecov
        uses: codecov/codecov-action@v3

  build-and-deploy:
    needs: test
    runs-on: ubuntu-latest
    if: github.ref == 'refs/heads/main'

    steps:
```

- **name:** Checkout code
uses: actions/checkout@v3

- **name:** Configure AWS credentials
uses: aws-actions/configure-aws-credentials@v2
with:
 - aws-access-key-id:** \${ secrets.AWS_ACCESS_KEY_ID }
 - aws-secret-access-key:** \${ secrets.AWS_SECRET_ACCESS_KEY }
 - aws-region:** \${ env.AWS_REGION }

- **name:** Login to Amazon ECR
id: login-ecr
uses: aws-actions/amazon-ecr-login@v1

- **name:** Build, tag, and push image to Amazon ECR
id: build-image
env:
 - ECR_REGISTRY:** \${ steps.login-ecr.outputs.registry }
 - IMAGE_TAG:** \${ github.sha }**run:** |


```
docker build -t $ECR_REGISTRY/$ECR_REPOSITORY:$IMAGE_TAG .
docker push $ECR_REGISTRY/$ECR_REPOSITORY:$IMAGE_TAG
echo "image=$ECR_REGISTRY/$ECR_REPOSITORY:$IMAGE_TAG" >> $GITHUB_OUTPUT
```

- **name:** Fill in the new image ID in the Amazon ECS task definition
id: task-def
uses: aws-actions/amazon-ecs-render-task-definition@v1
with:
 - task-definition:** \${ env.ECS_TASK_DEFINITION }
 - container-name:** flask-app
 - image:** \${ steps.build-image.outputs.image }

- **name:** Deploy Amazon ECS task definition
uses: aws-actions/amazon-ecs-deploy-task-definition@v1
with:
 - task-definition:** \${ steps.task-def.outputs.task-definition }
 - service:** \${ env.ECS_SERVICE }
 - cluster:** \${ env.ECS_CLUSTER }
 - wait-for-service-stability:** true

Blue-Green Deployment Strategy

Step 1: Setup Blue-Green Environment

bash

```
# Create two identical environments  
# Blue environment (current production)
```

```
aws ecs create-service \  
  --cluster flask-app-cluster \  
  --service-name flask-app-blue \  
  --task-definition flask-app-task:1 \  
  --desired-count 2 \  
  --launch-type FARGATE
```

```
# Green environment (new version)
```

```
aws ecs create-service \  
  --cluster flask-app-cluster \  
  --service-name flask-app-green \  
  --task-definition flask-app-task:2 \  
  --desired-count 2 \  
  --launch-type FARGATE
```

Step 2: Traffic Switching Script

bash

```
#!/bin/bash
```

```
# blue-green-deploy.sh
```

```
GREEN_TARGET_GROUP_ARN="arn:aws:elasticloadbalancing:us-east-1:123456789012:targetgroup/flask-a
```

```
BLUE_TARGET_GROUP_ARN="arn:aws:elasticloadbalancing:us-east-1:123456789012:targetgroup/flask-ap
```

```
LISTENER_ARN="arn:aws:elasticloadbalancing:us-east-1:123456789012:listener/app/flask-app-alb/12
```

```
echo "Starting Blue-Green Deployment..."
```

```
# Health check green environment
```

```
echo "Checking green environment health..."
```

```
GREEN_HEALTH=$(aws elbv2 describe-target-health --target-group-arn $GREEN_TARGET_GROUP_ARN --qu
```

```
if [ "$GREEN_HEALTH" = "healthy" ]; then
```

```
    echo "Green environment is healthy. Switching traffic..."
```

```
    # Switch traffic to green
```

```
    aws elbv2 modify-listener \
```

```
        --listener-arn $LISTENER_ARN \
```

```
        --default-actions Type=forward,TargetGroupArn=$GREEN_TARGET_GROUP_ARN
```

```
    echo "Traffic switched to green environment"
```

```
    # Wait and verify
```

```
    sleep 30
```

```
    # Scale down blue environment
```

```
    aws ecs update-service \
```

```
        --cluster flask-app-cluster \
```

```
        --service flask-app-blue \
```

```
        --desired-count 0
```

```
    echo "Blue-Green deployment completed successfully"
```

```
else
```

```
    echo "Green environment is not healthy. Deployment aborted."
```

```
    exit 1
```

```
fi
```

Monitoring and Logging

CloudWatch Logs Setup

```
bash
```

```
# Create Log group
```

```
aws logs create-log-group --log-group-name /aws/ecs/flask-app
```

```
# Create Log stream
```

```
aws logs create-log-stream \  
    --log-group-name /aws/ecs/flask-app \  
    --log-stream-name flask-app-stream
```

Application Monitoring with CloudWatch


```

# Enhanced Flask app with monitoring
import boto3
import time
from flask import Flask, jsonify, request, g
from datetime import datetime

app = Flask(__name__)

# CloudWatch client
cloudwatch = boto3.client('cloudwatch', region_name='us-east-1')

@app.before_request
def before_request():
    g.start_time = time.time()

@app.after_request
def after_request(response):
    # Calculate response time
    response_time = (time.time() - g.start_time) * 1000

    # Send metrics to CloudWatch
    try:
        cloudwatch.put_metric_data(
            Namespace='FlaskApp/Performance',
            MetricData=[
                {
                    'MetricName': 'ResponseTime',
                    'Value': response_time,
                    'Unit': 'Milliseconds',
                    'Dimensions': [
                        {
                            'Name': 'Endpoint',
                            'Value': request.endpoint or 'unknown'
                        }
                    ]
                }
            ],
            {
                'MetricName': 'RequestCount',
                'Value': 1,
                'Unit': 'Count',
                'Dimensions': [
                    {
                        'Name': 'StatusCode',
                        'Value': str(response.status_code)
                    }
                ]
            }
        )
    except Exception as e:
        # Log the error (this would typically be done with a logger)
        print(f"Error sending metrics to CloudWatch: {e}")

```

```

        }
    ]
)
except Exception as e:
    app.logger.error(f"Failed to send metrics: {e}")

return response

@app.route('/metrics')
def metrics():
    """Custom metrics endpoint"""
    return jsonify({
        "custom_metrics": {
            "uptime": time.time() - app.start_time,
            "timestamp": datetime.now().isoformat()
        }
    })

# Initialize start time
app.start_time = time.time()

```

Health Check Endpoint

python

```
@app.route('/health')
def health_check():
    """Comprehensive health check"""
    health_status = {
        "status": "healthy",
        "timestamp": datetime.now().isoformat(),
        "version": os.environ.get('APP_VERSION', 'unknown'),
        "environment": os.environ.get('FLASK_ENV', 'production')
    }

    # Check database connection
    try:
        # Example database health check
        # db.session.execute('SELECT 1')
        health_status["database"] = "connected"
    except Exception as e:
        health_status["status"] = "unhealthy"
        health_status["database"] = f"error: {str(e)}"

    # Check Redis connection
    try:
        if redis_client:
            redis_client.ping()
            health_status["redis"] = "connected"
        else:
            health_status["redis"] = "not_configured"
    except Exception as e:
        health_status["status"] = "unhealthy"
        health_status["redis"] = f"error: {str(e)}"

    status_code = 200 if health_status["status"] == "healthy" else 503
    return jsonify(health_status), status_code
```

Security Best Practices

Dockerfile Security

dockerfile

Security-hardened Dockerfile

FROM python:3.9-slim

Security updates

RUN apt-get update && apt-get upgrade -y && \
 apt-get install -y --no-install-recommends \
 gcc && \
 rm -rf /var/lib/apt/lists/* && \
 apt-get clean

Create non-root user early

RUN groupadd -r appuser && useradd -r -g appuser appuser

Set secure working directory

WORKDIR /app

Copy and install dependencies as root

COPY requirements.txt .

RUN pip install --no-cache-dir --upgrade pip && \
 pip install --no-cache-dir -r requirements.txt

Copy application code

COPY . .

Change ownership to non-root user

RUN chown -R appuser:appuser /app

Switch to non-root user

USER appuser

Remove unnecessary packages

RUN pip uninstall -y pip setuptools

Security Labels

LABEL security.policy="restricted" \
 maintainer="your-team@company.com"

Health check

HEALTHCHECK --interval=30s --timeout=10s --start-period=5s --retries=3 \
 CMD python -c "import requests; requests.get('http://localhost:5000/health')" || exit 1

EXPOSE 5000

CMD ["gunicorn", "--bind", "0.0.0.0:5000", "--workers", "2", "--timeout", "30", "app:app"]

AWS Security Configuration

bash

Create IAM role for ECS tasks

```
aws iam create-role \  
  --role-name ecsTaskRole \  
  --assume-role-policy-document '{  
    "Version": "2012-10-17",  
    "Statement": [  
      {  
        "Effect": "Allow",  
        "Principal": {  
          "Service": "ecs-tasks.amazonaws.com"  
        },  
        "Action": "sts:AssumeRole"  
      }  
    ]  
  }'
```

Attach minimal permissions policy

```
aws iam attach-role-policy \  
  --role-name ecsTaskRole \  
  --policy-arn arn:aws:iam::aws:policy/CloudWatchLogsFullAccess
```

Create security group with minimal permissions

```
aws ec2 create-security-group \  
  --group-name flask-app-secure-sg \  
  --description "Secure security group for Flask app"
```

Allow only necessary ports

```
aws ec2 authorize-security-group-ingress \  
  --group-name flask-app-secure-sg \  
  --protocol tcp \  
  --port 80 \  
  --source-group sg-alb-security-group-id
```

No direct SSH access - use Session Manager instead

Backup and Disaster Recovery

Database Backup Strategy

bash

Automated RDS backup

```
aws rds create-db-snapshot \  
  --db-instance-identifier flask-app-db \  
  --db-snapshot-identifier flask-app-backup-$(date +%Y%m%d%H%M%S)
```

Cross-region backup

```
aws rds copy-db-snapshot \  
  --source-db-snapshot-identifier flask-app-backup-20231201120000 \  
  --target-db-snapshot-identifier flask-app-backup-20231201120000-dr \  
  --source-region us-east-1 \  
  --target-region us-west-2
```

Application Data Backup

bash

S3 backup script

#!/bin/bash

backup.sh

BACKUP_DATE=\$(date +%Y%m%d_%H%M%S)

BACKUP_BUCKET="flask-app-backups"

Create backup archive

```
tar -czf app-backup-$BACKUP_DATE.tar.gz \  
  /var/www/flask-app \  
  --exclude='*.pyc' \  
  --exclude='__pycache__' \  
  --exclude='.git'
```

Upload to S3

```
aws s3 cp app-backup-$BACKUP_DATE.tar.gz s3://$BACKUP_BUCKET/backups/
```

Set Lifecycle policy for cost optimization

```
aws s3api put-bucket-lifecycle-configuration \  
  --bucket $BACKUP_BUCKET \  
  --lifecycle-configuration '{  
    "Rules": [  
      {  
        "ID": "BackupLifecycle",  
        "Status": "Enabled",  
        "Filter": {"Prefix": "backups/"},  
        "Transitions": [  
          {  
            "Days": 30,  
            "StorageClass": "STANDARD_IA"  
          },  
          {  
            "Days": 90,  
            "StorageClass": "GLACIER"  
          }  
        ]  
      }  
    ]  
  }'
```

Clean up local backup

```
rm app-backup-$BACKUP_DATE.tar.gz
```

Cost Optimization

Resource Right-sizing

bash

Monitor instance utilization

```
aws cloudwatch get-metric-statistics \
  --namespace AWS/EC2 \
  --metric-name CPUUtilization \
  --dimensions Name=InstanceId,Value=i-1234567890abcdef0 \
  --start-time 2023-11-01T00:00:00Z \
  --end-time 2023-12-01T00:00:00Z \
  --period 3600 \
  --statistics Average
```

Use Spot instances for non-critical workloads

```
aws ec2 request-spot-instances \
  --instance-count 2 \
  --type "one-time" \
  --launch-specification '{
    "ImageId": "ami-0c02fb55956c7d316",
    "InstanceType": "t3.medium",
    "KeyName": "my-flask-key",
    "SecurityGroups": ["flask-app-sg"]
  }'
```

Container Optimization

dockerfile

Multi-stage build for smaller images

```
FROM python:3.9 as builder
WORKDIR /app
COPY requirements.txt .
RUN pip install --user -r requirements.txt

FROM python:3.9-slim
WORKDIR /app
COPY --from=builder /root/.local /root/.local
COPY . .
ENV PATH=/root/.local/bin:$PATH
CMD ["gunicorn", "--bind", "0.0.0.0:5000", "app:app"]
```

Conclusion

This comprehensive guide covered the complete journey from Linux fundamentals to deploying containerized Flask applications on AWS. Key takeaways include:

Best Practices Summary

- 1. **Security First:** Always use non-root users, minimal permissions, and security groups
- 2. **Monitoring:** Implement comprehensive logging and monitoring from day one
- 3. **Automation:** Use Infrastructure as Code and CI/CD pipelines
- 4. **Scalability:** Design for horizontal scaling with load balancers and auto-scaling
- 5. **Cost Optimization:** Right-size resources and use appropriate storage classes
- 6. **Disaster Recovery:** Implement backup strategies and multi-region deployments

Deployment Decision Matrix

Use Case	Recommended Approach
Simple Apps	Direct EC2 + Nginx
Scalable Apps	ECS Fargate + ALB
Quick Prototyping	Elastic Beanstalk
Microservices	ECS + Service Discovery
High Traffic	ECS + Auto Scaling + CloudFront
Development	Docker Compose locally

Next Steps

- 1. Implement Infrastructure as Code using AWS CDK or Terraform
- 2. Add comprehensive testing and security scanning to CI/CD
- 3. Explore container orchestration with Amazon EKS
- 4. Implement observability with AWS X-Ray and CloudWatch Insights
- 5. Consider serverless options with AWS Lambda for event-driven workloads